

# Day 8: Database Implementation – NoSQL & Vector Databases

Constraint respected: **No Docker is used** anywhere in this day's implementation.

This day focuses on **document databases**, **vector databases**, and **cloud object storage**, which together form the backbone of modern **AI-driven backend systems**.

---

## 1. MongoDB – Document Model & Validation

### MongoDB Document Model

- Schema-flexible (but schema-controlled via validation)
- JSON-like documents (BSON)
- Optimized for nested, evolving data

### Database & Collection Setup

```
use ai_app
```

### Collection with Validation Rules

```
db.createCollection("users", {  
  validator: {  
    $jsonSchema: {  
      bsonType: "object",  
      required: ["name", "email", "created_at"],  
      properties: {  
        name: {  
          bsonType: "string",  
          description: "must be a string and is required"  
        },  
        email: {  
          bsonType: "string",  
          pattern: "^.+@.+$",  
          description: "must be a valid email"  
        },  
        created_at: {  
          bsonType: "date"  
        }  
      }  
    }  
  }  
}
```

```
        }
    }
})
```

## Insert & Query

```
db.users.insertOne({
  name: "Ashutosh",
  email: "ashu@example.com",
  created_at: new Date()
})

// Query
db.users.find({ email: "ashu@example.com" })
```

✓ Checklist: MongoDB collections with validation rules

## 2. Vector Databases – Conceptual Overview

### Why Vector Databases

- Traditional DBs fail at **semantic similarity**
- Vectors encode meaning, not keywords

Use cases: - Semantic search - Recommendation engines - RAG (Retrieval Augmented Generation) - Chat memory & AI assistants

## 3. ChromaDB (Local, No Docker)

### Installation

```
pip install chromadb sentence-transformers
```

### Initialize ChromaDB

```
import chromadb
from chromadb.config import Settings

client = chromadb.Client(
    Settings(persist_directory="./chroma_data")
)
```

```
collection = client.get_or_create_collection(name="docs")
```

## Insert Sample Data

```
documents = [  
    "PostgreSQL is a relational database",  
    "MongoDB is a document-oriented NoSQL database",  
    "Vector databases enable semantic search",  
]  
  
ids = ["doc1", "doc2", "doc3"]  
  
collection.add(documents=documents, ids=ids)
```

✓ Checklist: Vector database configured with sample data

---

## 4. Embedding Generation Pipeline

### SentenceTransformer-Based Embeddings

```
from sentence_transformers import SentenceTransformer  
  
model = SentenceTransformer("all-MiniLM-L6-v2")  
  
def generate_embedding(text: str):  
    return model.encode(text).tolist()
```

### Store with Embeddings

```
embeddings = [generate_embedding(d) for d in documents]  
  
collection.add(  
    documents=documents,  
    embeddings=embeddings,  
    ids=ids  
)
```

✓ Checklist: Embedding generation pipeline implemented

---

## 5. Semantic Search with Similarity Threshold

### Query with Threshold

```
query = "Which database is NoSQL?"  
query_embedding = generate_embedding(query)  
  
results = collection.query(  
    query_embeddings=[query_embedding],  
    n_results=3  
)  
  
for doc, score in zip(results['documents'][0], results['distances'][0]):  
    if score < 0.4: # similarity threshold  
        print("MATCH:", doc, "score:", score)
```

### Similarity Notes

- Lower distance = higher similarity
- Threshold tuning is **domain-specific**

✓ Checklist: Semantic search with threshold working

---

## 6. Weaviate (Conceptual + Optional Local Setup)

### When to Use Weaviate

- Hybrid search (vector + filters)
- Large-scale production systems
- Multi-tenant vector workloads

*(ChromaDB is enough for local & interview-level mastery)*

---

## 7. Cloud Object Storage – S3 & GCS

### Why Object Storage

- Store embeddings metadata, PDFs, images
  - Cheap, scalable, durable
-

## 8. AWS S3 Integration (Signed URLs)

### Install SDK

```
pip install boto3
```

### Generate Signed Upload URL

```
import boto3

s3 = boto3.client("s3")

url = s3.generate_presigned_url(
    ClientMethod="put_object",
    Params={
        "Bucket": "my-ai-bucket",
        "Key": "uploads/file.txt"
    },
    ExpiresIn=3600
)

print(url)
```

✓ Checklist: S3 signed URL integration complete

---

## 9. GCP Cloud Storage (Signed URLs)

### Install SDK

```
pip install google-cloud-storage
```

### Signed URL Example

```
from google.cloud import storage

client = storage.Client()
bucket = client.bucket("my-gcs-bucket")
blob = bucket.blob("uploads/file.txt")

url = blob.generate_signed_url(expiration=3600, method="PUT")
print(url)
```

✓ Checklist: GCS signed URL integration complete

---

## Day 8 Completion Status

- MongoDB collections with validation rules
  - Vector DB (ChromaDB) configured & tested
  - Embedding generation pipeline implemented
  - Semantic search with similarity threshold
  - AWS S3 & GCP signed URL integration
- 

## Architectural Insight (Important)

Typical AI backend stack:

```
Client → API (FastAPI)
    → MongoDB (metadata)
    → Vector DB (semantic retrieval)
    → LLM
    → S3 / GCS (documents, files)
```

---

Next logical progression: **Day 9 - Backend Performance & Scaling** (Caching, async IO, batching, N+1 problems, load testing)

Or we can **merge Days 6–8 into a single AI-ready backend architecture.**