

Day 14: Kubernetes & Container Orchestration

This module builds **core Kubernetes fundamentals with production context**. The focus is not just YAML, but *why Kubernetes exists, how its control plane works, and how applications are safely deployed, scaled, and updated.*

1. Background: Why Kubernetes Exists

The Problem Kubernetes Solves

Before Kubernetes: - Manual container deployment - No self-healing - No standardized scaling - Fragile deployments

Kubernetes provides: - **Scheduling** of containers - **Self-healing** (restart, reschedule) - **Service discovery & load balancing** - **Declarative desired state**

Kubernetes is an **orchestrator**, not a container runtime.

2. Kubernetes Architecture (Core Concepts)

Control Plane Components

Component	Responsibility
API Server	Entry point for all operations
etcd	Distributed key-value store (cluster state)
Scheduler	Decides where Pods run
Controller Manager	Maintains desired state

Worker Node Components

Component	Responsibility
kubelet	Ensures containers are running
kube-proxy	Networking & service routing
Container Runtime	Runs containers (containerd)

3. Local Kubernetes Cluster Setup

Option 1: Minikube

```
minikube start  
kubectl get nodes
```

Option 2: Kind

```
kind create cluster  
kubectl cluster-info
```

✓ Checklist: Kubernetes cluster operational

4. Kubernetes Core Objects

Pod

- Smallest deployable unit
- One or more containers

Deployment

- Manages ReplicaSets
- Enables rolling updates & rollbacks

Service

- Stable network endpoint
- Load-balancing across Pods

Ingress

- HTTP/HTTPS routing
- External access

5. Deployment Manifest (Application)

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: demo-app
```

```
spec:  
  replicas: 2  
  selector:  
    matchLabels:  
      app: demo  
  template:  
    metadata:  
      labels:  
        app: demo  
    spec:  
      containers:  
      - name: demo  
        image: demo/app:1.0  
        ports:  
        - containerPort: 8000
```

Apply:

```
kubectl apply -f deployment.yaml  
kubectl get deployments
```

✓ Checklist: Deployment applied successfully

6. Service Manifest (Expose Application)

```
apiVersion: v1  
kind: Service  
metadata:  
  name: demo-service  
spec:  
  selector:  
    app: demo  
  ports:  
  - port: 80  
    targetPort: 8000  
  type: ClusterIP
```

Verify:

```
kubectl get svc
```

- ✓ Checklist: Service exposing application
-

7. Ingress (External Access)

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: demo-ingress
spec:
  rules:
    - host: demo.local
      http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: demo-service
                port:
                  number: 80
```

Enable ingress controller (Minikube):

```
minikube addons enable ingress
```

8. ConfigMaps & Secrets (Environment Management)

ConfigMap

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  APP_ENV: production
  LOG_LEVEL: info
```

Secret

```
apiVersion: v1
kind: Secret
metadata:
  name: app-secrets
type: Opaque
data:
  DB_PASSWORD: cGFzc3dvcmQ=
```

Mount as environment variables:

```
envFrom:
- configMapRef:
  name: app-config
- secretRef:
  name: app-secrets
```

✓ Checklist: ConfigMaps & Secrets configured

9. Horizontal Pod Autoscaling (HPA)

Why HPA

- Handle traffic spikes automatically
- Optimize cost

Metrics Server

```
kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
```

HPA Manifest

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: demo-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
```

```
  name: demo-app
  minReplicas: 2
  maxReplicas: 5
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 60
```

Verify:

```
kubectl get hpa
```

✓ Checklist: HPA scaling based on metrics

10. Rolling Updates (Zero Downtime Deployments)

How Rolling Updates Work

- Gradually replace old Pods
- Maintain service availability

Update Image

```
kubectl set image deployment/demo-app demo=demo/app:2.0
```

Monitor rollout:

```
kubectl rollout status deployment/demo-app
```

Rollback if needed:

```
kubectl rollout undo deployment/demo-app
```

✓ Checklist: Rolling update performed successfully

Day 14 Completion Status

- Local Kubernetes cluster running
 - Deployments applied
 - Services & Ingress configured
 - ConfigMaps & Secrets used
 - HPA scaling verified
 - Rolling updates completed
-

Key Interview Insight

Kubernetes is about **desired state management**. You describe *what you want* — Kubernetes figures out *how to keep it running*.

Strong Next Steps

- Kubernetes networking deep dive
- Helm charts
- Kubernetes security (RBAC, NetworkPolicies)
- GitOps with ArgoCD

Tell me what you want to continue with.